



# **LungXAI Project Review Guide**

**Complete Documentation for Project Defense & Viva**

---

## TABLE OF CONTENTS

- Project Overview
  - Problem Statement & Motivation
  - Complete Architecture
  - Technology Stack - What, Where & Why
  - Module-by-Module Breakdown
  - Key Algorithms & Techniques
  - Dataset Details
  - Training Process
  - Evaluation Metrics
  - Potential Questions & Answers
  - Quick Reference Cards
-

# 1. PROJECT OVERVIEW

## 1.1 Project Title

"Explainable AI for Multi-Class Lung Cancer Classification Using Deep Learning and RAG-Based Knowledge Retrieval"

## 1.2 One-Line Summary

A deep learning system that not only classifies lung CT scan images into cancer types but also provides visual explanations (WHERE the model looks) and textual explanations (WHY those regions matter) using Grad-CAM and RAG.

## 1.3 Key Innovation

Traditional AI	Our System
Only predicts class	Predicts + Explains
Black-box	Transparent
No medical context	Provides medical knowledge
Single output	Visual + Textual explanation

## 1.4 Classes Classified (5 Classes)

Class	Description	Typical Location
Adenocarcinoma	Most common lung cancer (~40%)	Peripheral/Outer regions
Squamous Cell Carcinoma	Second most common (~30%)	Central/Hilar regions
Large Cell Carcinoma	Fast-growing, aggressive	Peripheral, larger masses
Benign Cases	Non-cancerous lesions	Variable
Normal Cases	Healthy lung tissue	Clear patterns

## 2. PROBLEM STATEMENT & MOTIVATION

### 2.1 The Problem

- Lung cancer is the leading cause of cancer deaths worldwide
- Early detection is critical - 5-year survival rate is 56% if detected early vs 5% if detected late
- Radiologist shortage - Increasing CT scan volumes with limited experts
- AI trust issue - Doctors don't trust black-box AI predictions

### 2.2 Why Explainability Matters

[X] Without Explainability:

AI says: "Adenocarcinoma (92%)"  
Doctor thinks: "But why? I can't trust this."

[OK] With Explainability:

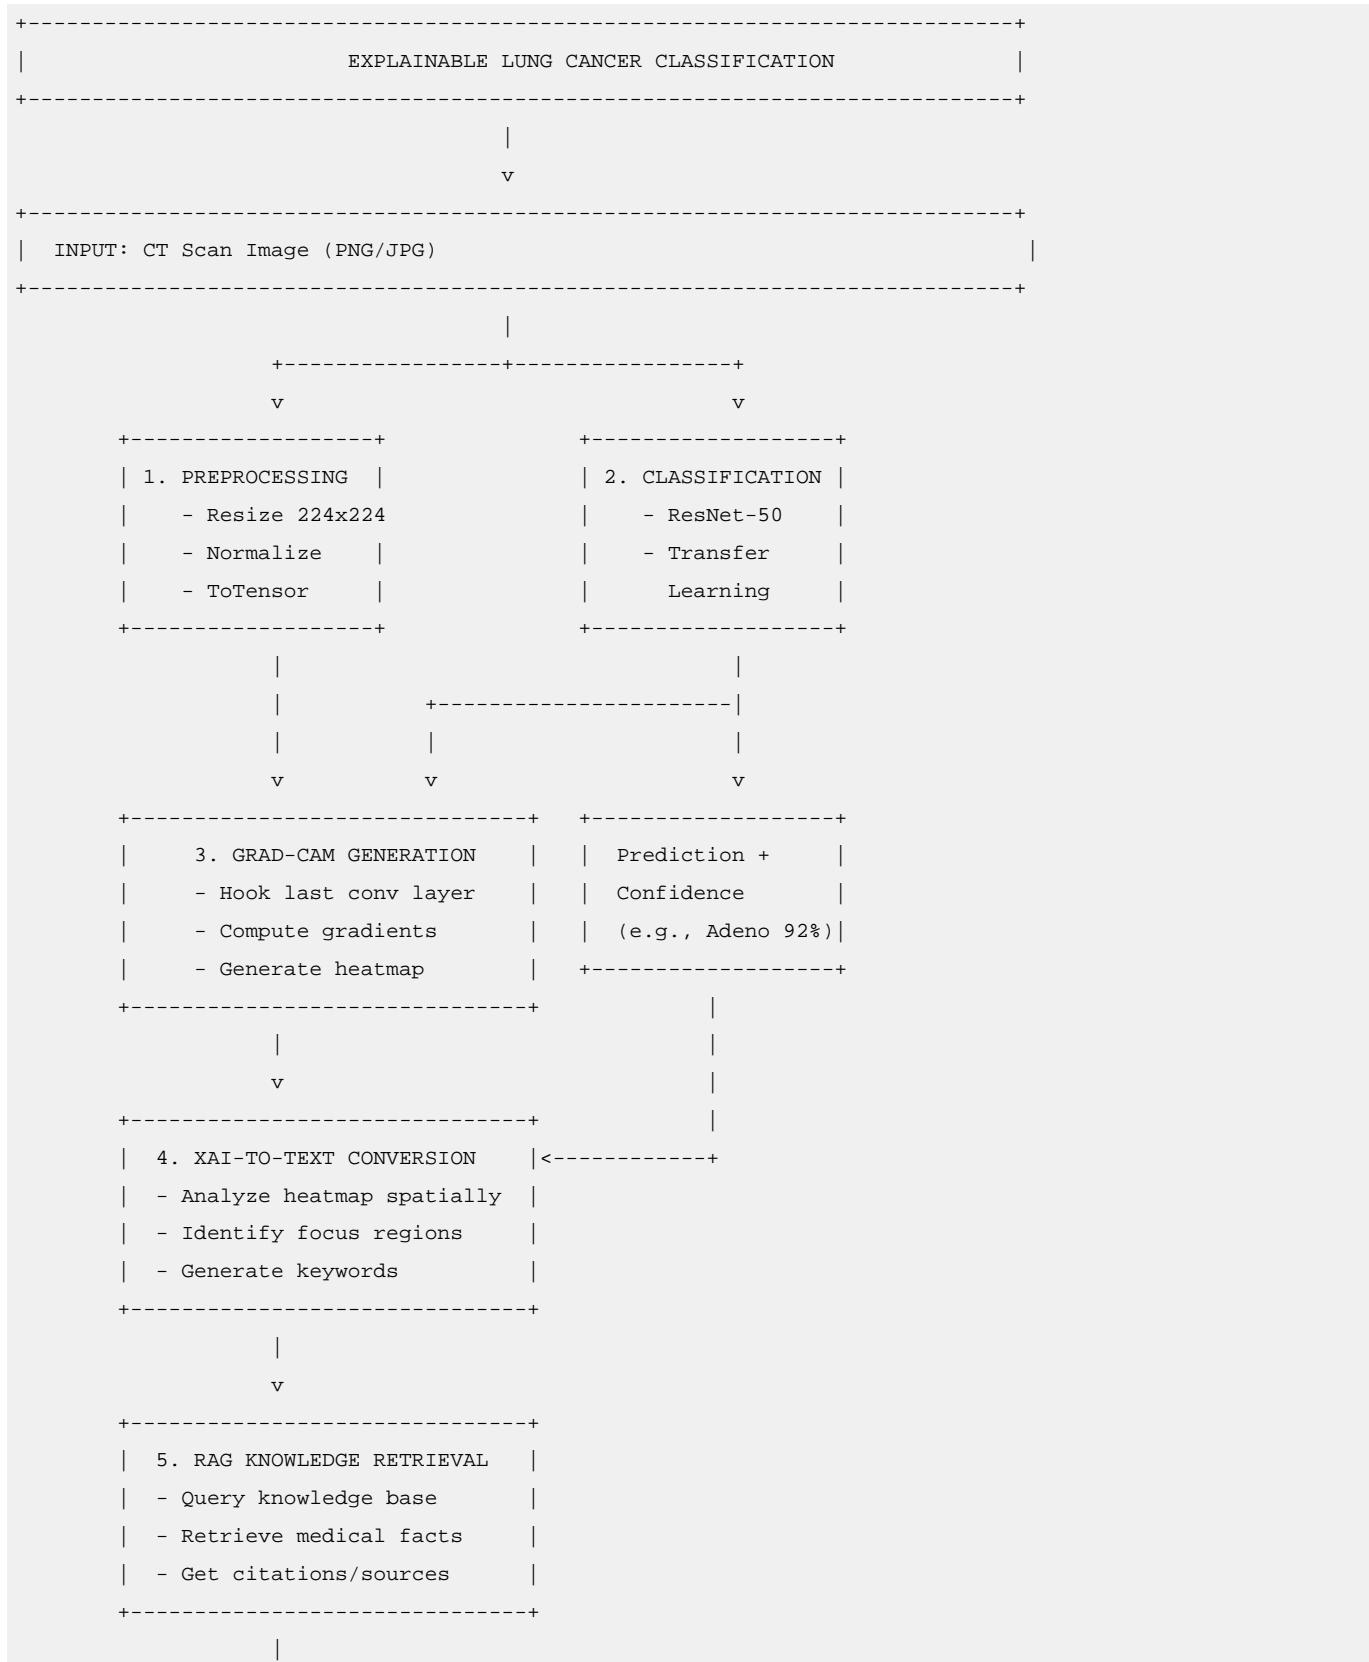
AI says: "Adenocarcinoma (92%)"  
+ Shows: Heatmap highlighting peripheral region  
+ Explains: "Adenocarcinoma typically presents in peripheral regions. Model focused on upper-right peripheral area with ground-glass opacity pattern."  
Doctor thinks: "This aligns with clinical knowledge. I can verify."

### 2.3 Project Objectives

- Classify lung CT images into 5 categories with high accuracy
- Visualize model attention using Grad-CAM
- Explain predictions using medical knowledge (RAG approach)
- Bridge the gap between AI predictions and clinical understanding

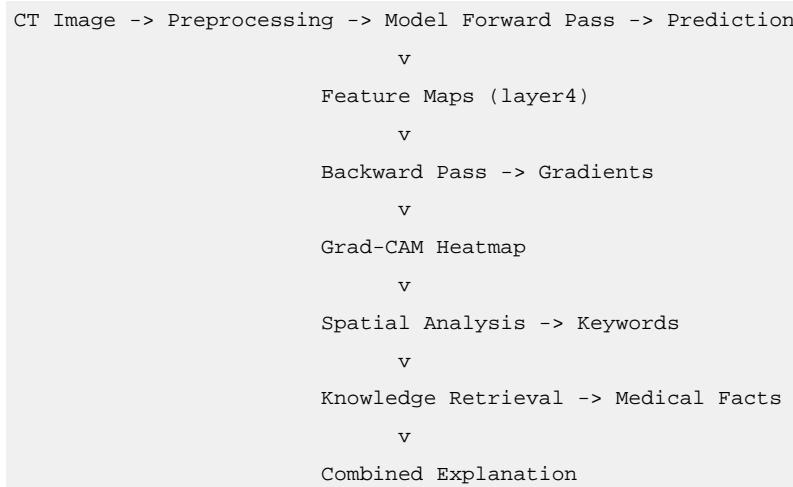
## 3. COMPLETE ARCHITECTURE

### 3.1 High-Level Pipeline



v			
OUTPUT:			
Original Image	Grad-CAM Heatmap (Where model focused)	Textual Explanation - Visual evidence - Medical context - Cited sources	

## 3.2 Data Flow Diagram



## 4. TECHNOLOGY STACK - WHAT, WHERE & WHY

### 4.1 Core Framework

Technology	What Is It	Where Used	Why Used
Python 3.10+	Programming language	Entire project	Industry standard for ML, e...
PyTorch 2.0+	Deep learning framework	Model, training, inference	Academic flexibility, dynam...
torchvision	Vision library	Pretrained models, transforms	Official PyTorch extension ...

### 4.2 Model Architecture

Technology	What Is It	Where Used	Why Used
ResNet-50	50-layer residual network	classifier.py	Proven performance, residua...
Transfer Learning	Reusing pretrained weights	Model initialization	ImageNet pretrained weights...
Dropout (0.5)	Regularization technique	Final FC layer	Prevents overfitting, impro...

### Why ResNet-50 Specifically?

ResNet-50 Benefits:

```
|-- [OK] 50 layers deep - captures complex patterns
|-- [OK] Residual connections - solves vanishing gradient
|-- [OK] 2048 features before FC - rich representations
|-- [OK] Well-documented - easy to explain in vivo
|-- [OK] Grad-CAM compatible - layer4 gives clear heatmaps
+-- [OK] ImageNet pretrained - 1000-class general features
```

Alternative Considered:

```
|-- VGG-16: Simpler but fewer features, no residuals
|-- EfficientNet: More efficient but harder to explain
+-- ViT: Newer but attention mechanism different from Grad-CAM
```

### 4.3 Explainability Stack

Technology	What Is It	Where Used	Why Used
Grad-CAM	Gradient-weighted Class Act...	gradcam.py	Visual explanation, shows i...
OpenCV	Computer vision library	visualize.py	Heatmap overlay, colormap a...
Matplotlib	Plotting library	Visualization	Creating publication-qualit...

## 4.4 RAG (Retrieval-Augmented Generation) Stack

Technology	What Is It	Where Used	Why Used
Custom Knowledge Base	JSON-structured medical facts	knowledge_base.py	Verified content, no halluc...
Keyword Matching	Simple retrieval method	Knowledge retrieval	Transparent, explainable, s...
XAI-to-Text Converter	Heatmap analysis	xai_to_text.py	Novel contribution - bridge...

## 4.5 Data Processing

Technology	What Is It	Where Used	Why Used
PIL (Pillow)	Image loading	dataset.py	Standard Python imaging lib...
NumPy	Numerical operations	Throughout	Array operations, statistics
scikit-learn	ML utilities	metrics.py, data splitting	Train/test split, evaluatio...
tqdm	Progress bars	Training loops	Visual progress tracking

## 4.6 Image Preprocessing Details

Transform	Parameter	Why
Resize	224x224	Standard input size for Ima...
Normalize	mean=[0.485,0.456,0.406], s...	ImageNet statistics - requi...
RandomHorizontalFlip	p=0.5	Data augmentation - lungs a...
RandomRotation	+/-15 deg	Data augmentation - handles...
ColorJitter	brightness/contrast=0.1	Data augmentation - handles...

### Why NOT Vertical Flip?

Anatomically, lungs have top-bottom orientation that matters clinically. Flipping vertically would create unrealistic images.

## 5. MODULE-BY-MODULE BREAKDOWN

### 5.1 src/data/ - Data Handling

#### dataset.py - Custom PyTorch Dataset

```
class LungCancerDataset(Dataset):
    """
    WHAT: Custom dataset class for loading CT images
    WHY: Full control over data loading, can add metadata later
    HOW:
        - Scans class folders
        - Loads images with PIL
        - Applies transforms
        - Returns (image_tensor, label) pairs
    """

```

Key Methods:

Method	Purpose
<code>__init__</code>	Initialize paths, detect cl...
<code>__len__</code>	Return total number of images
<code>__getitem__</code>	Load single image, apply tr...

#### transforms.py - Image Preprocessing

```
# Training: Augmentation + Normalization
get_train_transforms() -> Resize -> Flip -> Rotate -> ColorJitter -> ToTensor -> Normalize

# Validation/Test: Only Normalization (no augmentation)
get_val_transforms() -> Resize -> ToTensor -> Normalize
```

Why Separate Transforms?

*Training needs augmentation to increase data variety. Validation/Test needs deterministic, consistent preprocessing for reliable evaluation.*

#### dataloader.py - DataLoader Creation

```
# Creates three dataloaders with proper splitting
train_loader, val_loader, test_loader = create_dataloaders(...)

# Split ratios: 70% train, 15% val, 15% test
```

```
# Stratified split ensures each set has proportional class representation
```

---

## 5.2 src/models/ - Neural Network

### classifier.py - Main Model

```
class LungCancerClassifier(nn.Module):  
    """  
    Architecture:  
    +-----+  
    | ResNet-50 Backbone |  
    +-----+  
    | conv1 -> bn1 -> relu -> maxpool |  
    | v |  
    | layer1 (64 -> 256 channels) - 3 Bottleneck blocks |  
    | v |  
    | layer2 (256 -> 512 channels) - 4 Bottleneck blocks |  
    | v |  
    | layer3 (512 -> 1024 channels) - 6 Bottleneck blocks |  
    | v |  
    | layer4 (1024 -> 2048 channels) - 3 Bottleneck blocks <- Grad-CAM target |  
    | v |  
    | avgpool (adaptive) |  
    | v |  
    | Dropout(0.5) |  
    | v |  
    | fc (2048 -> 5 classes) - Modified for our task |  
    +-----+  
    """
```

### Key Implementation Details:

```
# Load pretrained weights  
weights = models.ResNet50_Weights.IMGNET1K_V2  
self.backbone = models.resnet50(weights=weights)  
  
# Modify final layer for our task  
# Original: 2048 -> 1000 (ImageNet)  
# Modified: 2048 -> 5 (Our classes)  
self.backbone.fc = nn.Sequential(  
    nn.Dropout(p=0.5),  
    nn.Linear(2048, 5)  
)
```

### model\_factory.py - Model Creation

```
def create_model(model_name, num_classes, pretrained, ...):
    """
    Factory pattern for creating models
    - Centralizes model creation
    - Easy to add new architectures
    - Consistent configuration
    """

```

---

## 5.3 src/xai/ - Explainable AI

### gradcam.py - Grad-CAM Implementation

What is Grad-CAM?

*Gradient-weighted Class Activation Mapping - uses gradients flowing into the final convolutional layer to produce a heatmap highlighting important regions.*

How It Works (Step by Step):

```
Step 1: Forward Pass
        Input image -> Model -> Feature maps (A) captured at layer4

Step 2: Get Target Score
        Output logits -> Select score for target class (yc)

Step 3: Backward Pass
        Compute dyc/dA (gradients of target class w.r.t. feature maps)

Step 4: Global Average Pooling
        alphak = (1/Z) Sumi Sumj (dyc/dAkij)
        (Importance weight for each feature map k)

Step 5: Weighted Combination
        L = ReLU(Sumk alphak ? Ak)
        (Combine feature maps weighted by importance, keep only positive)

Step 6: Upscale
        Resize heatmap to input image size (224x224)
```

Implementation:

```
class GradCAM:
    def __init__(self, model, target_layer):
        # Register hooks to capture activations and gradients
        self.forward_handle = target_layer.register_forward_hook(forward_hook)
        self.backward_handle = target_layer.register_full_backward_hook(backward_hook)
```

```

def generate(self, input_tensor, target_class=None):
    # 1. Forward pass
    output = self.model(input_tensor)

    # 2. Get target class score
    target_score = output[0, target_class]

    # 3. Backward pass
    target_score.backward()

    # 4. Compute weights (global average pooling of gradients)
    weights = self.gradients.mean(dim=[2, 3], keepdim=True)

    # 5. Weighted combination with ReLU
    cam = (weights * self.activations).sum(dim=1)
    cam = F.relu(cam)

    # 6. Normalize to [0, 1]
    cam = (cam - cam.min()) / (cam.max() - cam.min())

    return cam

```

## Why layer4?

*Layer4 is the last convolutional layer in ResNet-50. It has:*

- Highest semantic meaning (understands "what" is in the image)
- 7x7 spatial resolution (enough for localization)
- 2048 channels (rich feature representations)

## visualize.py - Visualization Utilities

```

def create_heatmap_overlay(image, heatmap, alpha=0.4):
    """
    Overlays colored heatmap on original image

    Process:
    1. Normalize heatmap to [0, 255]
    2. Apply JET colormap (blue->green->yellow->red)
    3. Blend with original image (alpha blending)
    """

def visualize_gradcam(image, heatmap, predicted_class, confidence):
    """
    Creates 3-panel visualization:
    [Original Image] [Heatmap Only] [Overlay]
    """

```

## 5.4 src/rag/ - RAG Pipeline

### THIS IS THE NOVEL CONTRIBUTION

#### knowledge\_base.py - Medical Knowledge Store

```
class MedicalKnowledgeBase:  
    """  
    WHAT: Structured repository of medical knowledge  
  
    WHY Local (not LLM)?  
    |-- [OK] Controlled content - all facts are verified  
    |-- [OK] No hallucination - deterministic responses  
    |-- [OK] Citable - every fact has a source  
    |-- [OK] Works offline - no API dependency  
    +-- [OK] Explainable - can show exactly what knowledge is used  
  
    STRUCTURE:  
    {  
        "id": "adeno_001",  
        "keywords": ["adenocarcinoma", "peripheral", "outer"],  
        "content": "Adenocarcinoma typically presents in peripheral...",  
        "source": "Travis WD et al., WHO Classification, 2021"  
    }  
    """
```

#### Knowledge Retrieval Process:

```
Query: "adenocarcinoma peripheral upper"  
      v  
Tokenize query into keywords  
      v  
Search keyword index  
      v  
Rank by keyword matches  
      v  
Return top-k entries with sources
```

#### xai\_to\_text.py - XAI-to-Text Conversion (NOVEL!)

```
class XAITextConverter:  
    """  
    THE BRIDGE FROM VISUAL TO TEXTUAL  
  
    Input: Grad-CAM heatmap [H, W]  
    Output: Textual description + Keywords for RAG
```

```
Process:
```

1. Analyze spatial distribution (divide into 3x3 grid)
  2. Find high-attention regions
  3. Determine location (peripheral/central, upper/lower)
  4. Quantify intensity (very\_high, high, moderate, low)
  5. Generate natural language description
  6. Extract keywords for knowledge base query
- """

### Spatial Analysis:

top_left	top_center	top_right
(upper   left)	(upper   central)	(upper   right)
<hr/>		
middle_left	middle_center	middle_right
(middle   left)	(central)	(middle   right)
<hr/>		
bottom_left	bottom_center	bottom_right
(lower   left)	(lower   central)	(lower   right)
<hr/>		

For each region: Calculate mean attention value

Determine: Which regions have highest attention?

Map to: Medical terminology (peripheral = outer regions, central = middle)

## explanation\_generator.py - Final Explanation

```
class ExplanationGenerator:
    """
    Combines everything into coherent explanation

    Components:
    1. XAITextConverter - "Model focused on peripheral upper region"
    2. MedicalKnowledgeBase - "Adenocarcinoma typically presents peripherally"

    Output Format:
```

PREDICTION: Adenocarcinoma (92.3%)	
<hr/>	
VISUAL EVIDENCE:	
The model focused on the peripheral upper region with	
high attention intensity, concentrated in specific areas.	
<hr/>	
MEDICAL CONTEXT:	
Adenocarcinoma typically presents in the peripheral regions	
of the lung. Ground-glass opacity is frequently associated	

```
| with this cancer type. |
|-----|
| Sources: Travis WD et al., WHO Classification 2021 |
+-----+
'''
```

## 5.5 src/utils/ - Utilities

### config.py - Centralized Configuration

```
@dataclass
class Config:
    # Paths
    base_dir = "d:\\Major Project"
    dataset_dir = "archive (1)/Lung Cancer Dataset"

    # Classes
    class_names = ["adenocarcinoma", "Benign cases",
                   "large cell carcinoma", "Normal cases",
                   "squamous cell carcinoma"]
    num_classes = 5

    # Training
    image_size = (224, 224)
    batch_size = 32
    num_epochs = 10
    learning_rate = 1e-4

    # Model
    model_name = "resnet50"
    pretrained = True
    dropout_rate = 0.5
```

### metrics.py - Evaluation Metrics

```
# Metrics used:
accuracy_score()    # Overall correctness
precision_score()  # How many predicted positives are true
recall_score()     # How many actual positives were found
f1_score()         # Harmonic mean of precision & recall
confusion_matrix() # Full breakdown of predictions
```

Why Recall is Important in Medical Imaging:

*Missing a cancer case (False Negative) is more dangerous than a false alarm (False Positive). We want*

*HIGH RECALL to catch all potential cancers.*

---

## 6. KEY ALGORITHMS & TECHNIQUES

### 6.1 Transfer Learning

Concept:

```
ImageNet (1.2M images, 1000 classes)
    v Pre-training
    ResNet-50 learns general features
    (edges, textures, shapes, objects)
    v Transfer
    Load these weights
    v Fine-tune
    Train on our lung cancer dataset
    (Only modify final layers significantly)
```

Why It Works:

*Low-level features (edges, textures) learned on natural images transfer well to medical images. High-level features get adapted during fine-tuning.*

### 6.2 Residual Connections (ResNet)

Problem: Very deep networks suffer from vanishing gradients

Solution: Skip connections that add input to output

```
Traditional:      Residual:
x              x -----
v              v   |
Conv          Conv   |
v              v   |
Conv          Conv   |
v              v   |
y              y = F(x) + x -----
                  +--- Residual connection
```

Mathematical Form:

$$H(x) = F(x) + x$$

Where:

- $x$  is the input
- $F(x)$  is the learned residual
- $H(x)$  is the output

## 6.3 Cross-Entropy Loss

Used for: Multi-class classification

Formula:

$$L = -\sum_{c=1}^C y_c \log(p_c)$$

Where:

- C = number of classes (5)
- $y_c$  = 1 if correct class, 0 otherwise
- $p_c$  = predicted probability for class c

## 6.4 AdamW Optimizer

Why AdamW?

- Adaptive learning rate per parameter
- Weight decay for regularization
- Works well with pretrained models

Key Parameters:

```
optimizer = torch.optim.AdamW(  
    model.parameters(),  
    lr=1e-4,           # Learning rate  
    weight_decay=1e-4  # L2 regularization  
)
```

---

## 7. DATASET DETAILS

### 7.1 Dataset Source

Kaggle Lung Cancer CT Scan Dataset

### 7.2 Dataset Structure

```
archive (1)/Lung Cancer Dataset/
|-- adenocarcinoma/      # ~1500 images
|-- Benign cases/        # ~1500 images
|-- large cell carcinoma/ # ~1500 images
|-- Normal cases/        # ~1500 images
+-- squamous cell carcinoma/ # ~1500 images
```

### 7.3 Image Characteristics

Property	Value
Format	PNG/JPG
Original Size	Variable
Processed Size	224x224
Color	RGB (converted if grayscale)
Source	CT scan slices

### 7.4 Data Split

Split	Percentage	Purpose
Training	70%	Learning model weights
Validation	15%	Hyperparameter tuning, earl...
Test	15%	Final unbiased evaluation

Stratified Split: Ensures each split has proportional class representation

## 8. TRAINING PROCESS

### 8.1 Training Pipeline

```

for epoch in range(num_epochs):
    # Training Phase
    model.train()
    for batch in train_loader:
        1. Forward pass: predictions = model(images)
        2. Compute loss: loss = criterion(predictions, labels)
        3. Backward pass: loss.backward()
        4. Update weights: optimizer.step()

    # Validation Phase
    model.eval()
    with torch.no_grad():
        Evaluate on validation set
        Track validation accuracy

    # Save best model if validation improves
    if val_accuracy > best_accuracy:
        save_checkpoint(model, "best_model.pth")

```

### 8.2 Training Configuration

Hyperparameter	Value	Reason
Epochs	10	Sufficient for convergence ...
Batch Size	32	Balance between memory and ...
Learning Rate	1e-4	Conservative for fine-tunin...
Optimizer	AdamW	Adaptive + weight decay
Loss	CrossEntropyLoss	Standard for multi-class

### 8.3 Checkpoints Saved

File	Description
best_model.pth	Model with highest validation accuracy
final_model.pth	Model after all epochs completed

## 9. EVALUATION METRICS

### 9.1 Primary Metrics

#### Accuracy

$$\text{Accuracy} = \frac{\text{Correct Predictions}}{\text{Total Predictions}}$$

#### Precision (Per Class)

$$\text{Precision} = \frac{\text{TP}}{\text{TP} + \text{FP}}$$

"Of all predicted as X, how many are actually X?"

#### Recall (Per Class)

$$\text{Recall} = \frac{\text{TP}}{\text{TP} + \text{FN}}$$

"Of all actual X, how many did we find?"

#### F1-Score

$$F1 = 2 \times \frac{\text{Precision} \times \text{Recall}}{\text{Precision} + \text{Recall}}$$

Harmonic mean of precision and recall

### 9.2 Confusion Matrix

		Predicted				
		A	SCC	LCC	B	N
Actual	A	[ TP	...	...	...	...
	SCC	[ ...	TP	...	...	...
	LCC	[ ...	...	TP	...	...
	B	[ ...	...	...	TP	...
	N	[ ...	...	...	...	TP ]

Where: A=Adenocarcinoma, SCC=Squamous, LCC=Large Cell, B=Benign, N=Normal

## 10. POTENTIAL QUESTIONS & ANSWERS

### Architecture Questions

#### Q1: Why ResNet-50 and not other architectures?

A: ResNet-50 offers the best balance for our use case:

- Proven medical imaging performance - widely used in literature
- Residual connections - solve vanishing gradient in deep networks
- 2048 features - rich representations for cancer patterns
- Grad-CAM compatible - layer4 produces clear, interpretable heatmaps
- Well-documented - extensive research papers and tutorials
- Transfer learning ready - ImageNet pretraining provides good initialization

#### Q2: Why 224x224 image size?

A: This is the standard input size for ImageNet-pretrained models:

- The pretrained convolutional filters expect this resolution
- Changing size would require architecture modifications
- It provides sufficient detail for classification while being computationally efficient

#### Q3: What are residual connections and why do they matter?

A: Residual connections add the input directly to the output:  $H(x) = F(x) + x$

- Solves vanishing gradient: Gradients can flow directly through skip connections
- Enables depth: Without them, networks >20 layers become hard to train
- Identity mapping: If  $F(x)=0$ , the block simply passes input through

#### Q4: Why transfer learning for medical imaging?

A: Medical datasets are typically small (thousands vs millions). Transfer learning helps because:

- Feature reuse: Low-level features (edges, textures) transfer across domains
- Better initialization: Pretrained weights start closer to optimal
- Faster convergence: Fewer epochs needed
- Reduced overfitting: More generalizable features

### Explainability Questions

## Q5: What is Grad-CAM and how does it work?

A: Grad-CAM (Gradient-weighted Class Activation Mapping) creates visual explanations:

- Forward pass: Get feature maps from target layer (layer4)
- Backward pass: Compute gradients of target class score w.r.t. feature maps
- Weight computation: Global average pool gradients to get importance weights
- Combine: Weighted sum of feature maps
- ReLU: Keep only positive activations (regions that help the prediction)
- Resize: Upscale to original image size

## Q6: Why Grad-CAM instead of other XAI methods?

A:

Method	Pros	Cons	Why Not
Grad-CAM	No modification needed, cla...	Coarse resolution	[OK] Selected
LIME	Model-agnostic	Slow, many parameters	Too complex
SHAP	Theoretically sound	Computationally expensive	Too slow
Attention	Built-in for transformers	Requires architecture change	Not applicable to ResNet

## Q7: What is the novel contribution - XAI-to-Text?

A: The XAI-to-Text converter bridges visual explanations with textual knowledge:

- Analyzes heatmap spatially - divides into 3x3 grid
- Identifies attention regions - peripheral, central, upper, lower
- Quantifies intensity - very high, high, moderate, low
- Generates keywords - "adenocarcinoma peripheral upper"
- Queries knowledge base - retrieves relevant medical facts

This is novel because it connects WHERE the model looks with WHY that matters medically.

## RAG Questions

## Q8: What is RAG and why use it?

A: RAG = Retrieval-Augmented Generation

- Concept: Instead of generating explanations from scratch, retrieve relevant facts from a curated knowledge base
- Why: Ensures accuracy, provides citations, no hallucination
- Implementation: Simple keyword matching (can upgrade to semantic search later)

## Q9: Why not use an LLM like GPT for explanations?

A:

LLM Approach	Our Approach
Can hallucinate	Deterministic, verified facts
Black-box generation	Transparent retrieval
Needs API, costs money	Works offline, free
Hard to cite sources	Every fact has a citation
May be wrong	Content is curated by us

For medical applications, reliability and citation are critical.

## Q10: How is the knowledge base structured?

A: JSON structure with:

```
{
  "id": "adeno_001",
  "keywords": [ "adenocarcinoma", "peripheral", "outer" ],
  "content": "Medical fact about adenocarcinoma...",
  "source": "Travis WD et al., WHO Classification, 2021"
}
```

- Keywords: Trigger terms for retrieval
- Content: Verified medical information
- Source: Academic citation

## Training Questions

### Q11: Why 70-15-15 split?

A: Standard split for ML projects:

- 70% Training: Main learning data
- 15% Validation: Tune hyperparameters, monitor for overfitting
- 15% Test: Final unbiased evaluation (never seen during training)

Stratified split ensures each set has proportional class distribution.

### Q12: What is data augmentation and why use it?

A: Data augmentation artificially increases training data variety:

- Horizontal flip: Lungs are roughly symmetric

- Rotation (+/-15 deg): Handles slight orientation variations
- Color jitter: Handles CT scanner variations

Why: Prevents overfitting, improves generalization, especially important with limited medical data.

### **Q13: Why use dropout?**

A: Dropout randomly sets neurons to 0 during training:

- Rate 0.5: 50% of neurons dropped each iteration
- Effect: Prevents co-adaptation of neurons
- Result: More robust features, better generalization

### **Q14: Why AdamW optimizer?**

A: AdamW = Adam with decoupled weight decay

- Adaptive learning rate: Different rate per parameter
- Momentum: Smooths gradient updates
- Weight decay: L2 regularization for better generalization
- Works well with pretrained models

## **Metrics Questions**

### **Q15: Why is recall important in medical imaging?**

A: In cancer detection:

- False Negative (missing cancer) = Dangerous, patient doesn't get treatment
- False Positive (false alarm) = Less dangerous, just additional tests

We want HIGH RECALL to minimize false negatives.

### **Q16: What does the confusion matrix tell us?**

A: Shows detailed breakdown:

- Diagonal = correct predictions
- Off-diagonal = where model confuses classes
- Helps identify: Which cancers are confused with each other?

## **Technical Questions**

## Q17: What are forward and backward hooks?

A: PyTorch hooks capture intermediate values:

- Forward hook: Captures output of a layer during forward pass
- Backward hook: Captures gradients during backward pass
- Usage in Grad-CAM: Capture feature maps and gradients for heatmap computation

## Q18: What is the model's output format?

A:

- Raw output: Logits (unnormalized scores) for each class
- After softmax: Probabilities summing to 1
- Prediction: argmax of probabilities

## Q19: How do you handle class imbalance?

A:

- Stratified split: Ensures proportional class distribution
- Weighted loss (optional): Give more weight to minority classes
- Data augmentation: Increases effective training data
- Metrics: Use macro-averaged F1 instead of accuracy

## Q20: What happens during inference (prediction)?

A:

1. Load image -> Preprocess (resize, normalize)
2. Forward pass -> Get logits
3. Softmax -> Get probabilities
4. Grad-CAM -> Get heatmap
5. XAI-to-Text -> Get description
6. RAG -> Get medical context
7. Combine -> Final explanation

# 11. QUICK REFERENCE CARDS

## 11.1 Pipeline Summary Card

```
+-----+
|          LUNGXAI PIPELINE SUMMARY      |
+-----+
| INPUT:  CT Scan Image (any size, PNG/JPG)    |
|      v                                         |
| PREPROCESS: Resize 224x224, Normalize        |
|      v                                         |
| MODEL: ResNet-50 (pretrained, fine-tuned)     |
|      v                                         |
| OUTPUT: 5-class probabilities                 |
|      v                                         |
| XAI: Grad-CAM heatmap (layer4)               |
|      v                                         |
| CONVERT: XAI-to-Text (spatial analysis)       |
|      v                                         |
| RAG: Knowledge retrieval (keyword matching)   |
|      v                                         |
| FINAL: Visual + Textual explanation          |
+-----+
```

## 11.2 Technology Quick Reference

Component	Technology	File
Model	ResNet-50 + PyTorch	classifier.py
XAI	Grad-CAM	gradcam.py
Knowledge	JSON Knowledge Base	knowledge_base.py
Bridge	XAI-to-Text Converter	xai_to_text.py
Visualization	OpenCV + Matplotlib	visualize.py

## 11.3 Key Numbers

Metric	Value
Input Size	224 x 224 x 3
Model Parameters	~25.6M (ResNet-50)
Feature Size	2048 (before FC)
Output Classes	5
Grad-CAM Layer	layer4 (7x7x2048)

Dropout Rate	0.5
Learning Rate	1e-4

## 11.4 Class Information Card

Class	Typical CT Features	Typical Location
Adenocarcinoma	Ground-glass opacity, spicules	Peripheral
Squamous Cell	Central mass, cavitation	Central/Hilar
Large Cell	Large peripheral mass	Peripheral
Benign	Well-defined margins, calcifications	Variable
Normal	Clear lung fields	N/A

## 11.5 Formula Quick Reference

Grad-CAM:  $L = \text{ReLU}(\sum_k \alpha_k \cdot A_k)$  where  $\alpha_k = \text{GAP}(dyc/dA_k)$

Loss:  $L = -\sum y_c \log(p_c)$

Accuracy:  $TP + TN / \text{Total}$

Precision:  $TP / (TP + FP)$

Recall:  $TP / (TP + FN)$

F1:  $2 \times (P \times R) / (P + R)$

## FINAL TIPS FOR REVIEW

- Start with the big picture - Explain the problem, then the solution
- Know the "why" - Be prepared to justify every technology choice
- Understand the flow - Be able to trace data through the entire pipeline
- Remember key numbers - 224x224, 5 classes, 2048 features, layer4
- Emphasize novelty - XAI-to-Text conversion is your unique contribution
- Clinical relevance - Always connect back to why this helps doctors

Good luck with your review!

---

Document generated for Major Project Review - January 2026